

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df = pd.read_csv('diabetes-data.csv')
```

```
In [3]: 1 df
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns



```
In [4]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction              768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [5]: 1 df.shape
```

```
Out[5]: (768, 9)
```

In [6]: 1 df.describe()

Out[6]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diab
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

In [7]: 1 df\_copy = df.copy(deep=True)

In [8]: 1 df\_copy

Out[8]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns

In [9]: 1 df\_copy[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] =

In [10]: 1 df\_copy

Out[10]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148.0	72.0	35.0	NaN	33.6	
1	1	85.0	66.0	29.0	NaN	26.6	
2	8	183.0	64.0	NaN	NaN	23.3	
3	1	89.0	66.0	23.0	94.0	28.1	
4	0	137.0	40.0	35.0	168.0	43.1	
...	...	...	...	...	...	...	
763	10	101.0	76.0	48.0	180.0	32.9	
764	2	122.0	70.0	27.0	NaN	36.8	
765	5	121.0	72.0	23.0	112.0	26.2	
766	1	126.0	60.0	NaN	NaN	30.1	
767	1	93.0	70.0	31.0	NaN	30.4	

768 rows × 9 columns

In [11]: 1 df\_copy.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              763 non-null    float64
2   BloodPressure                        733 non-null    float64
3   SkinThickness                        541 non-null    float64
4   Insulin                              394 non-null    float64
5   BMI                                  757 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(6), int64(3)
memory usage: 54.1 KB
```

In [12]: 1 df\_copy.isnull().sum()

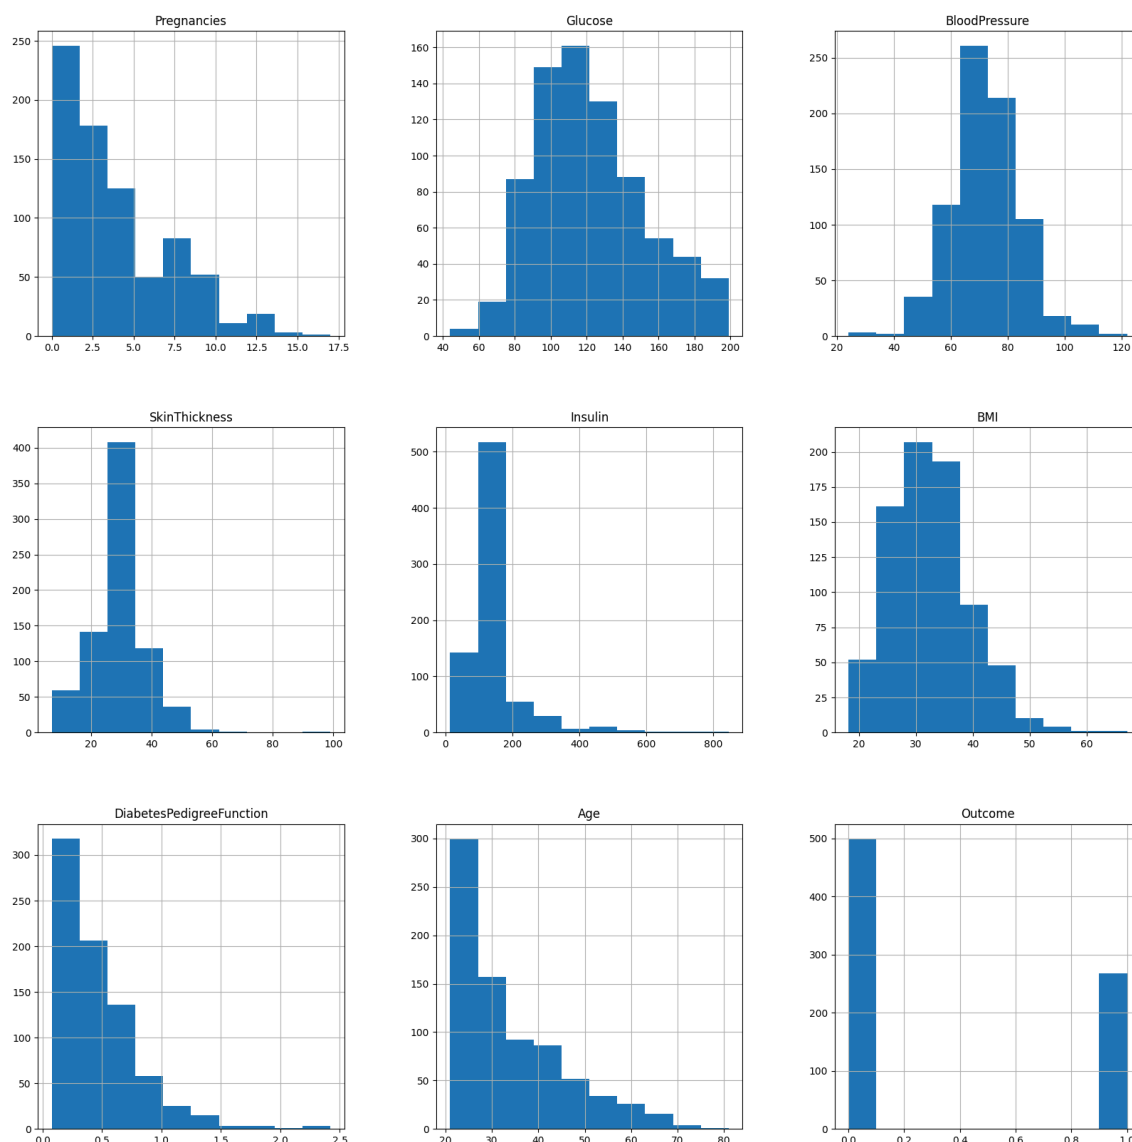
Out[12]:

Pregnancies	0
Glucose	5
BloodPressure	35
SkinThickness	227
Insulin	374
BMI	11
DiabetesPedigreeFunction	0
Age	0
Outcome	0
dtype: int64	

```
In [13]: 1 hplot = df.hist(figsize=(20,20))
```

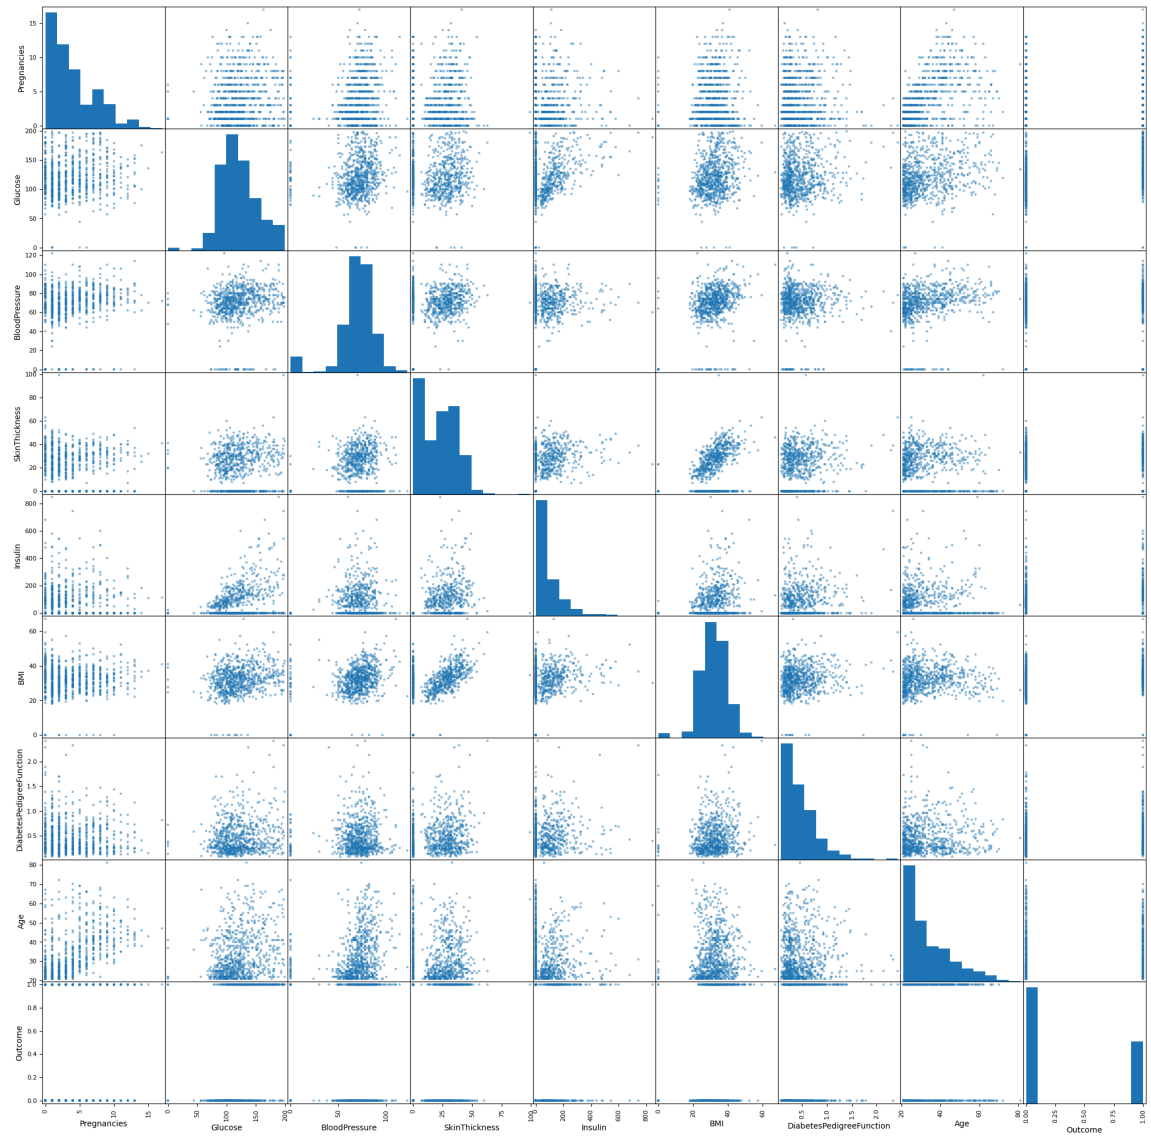


```
In [14]: 1 df_copy['Glucose'].fillna(df_copy['Glucose'].mean(),inplace=True)
2 df_copy['BloodPressure'].fillna(df_copy['BloodPressure'].mean(),inplace=True)
3 df_copy['SkinThickness'].fillna(df_copy['SkinThickness'].median(),inplace=True)
4 df_copy['Insulin'].fillna(df_copy['Insulin'].median(),inplace=True)
5 df_copy['BMI'].fillna(df_copy['BMI'].median(),inplace=True)
6 hplot=df_copy.hist(figsize=(20,20))
```

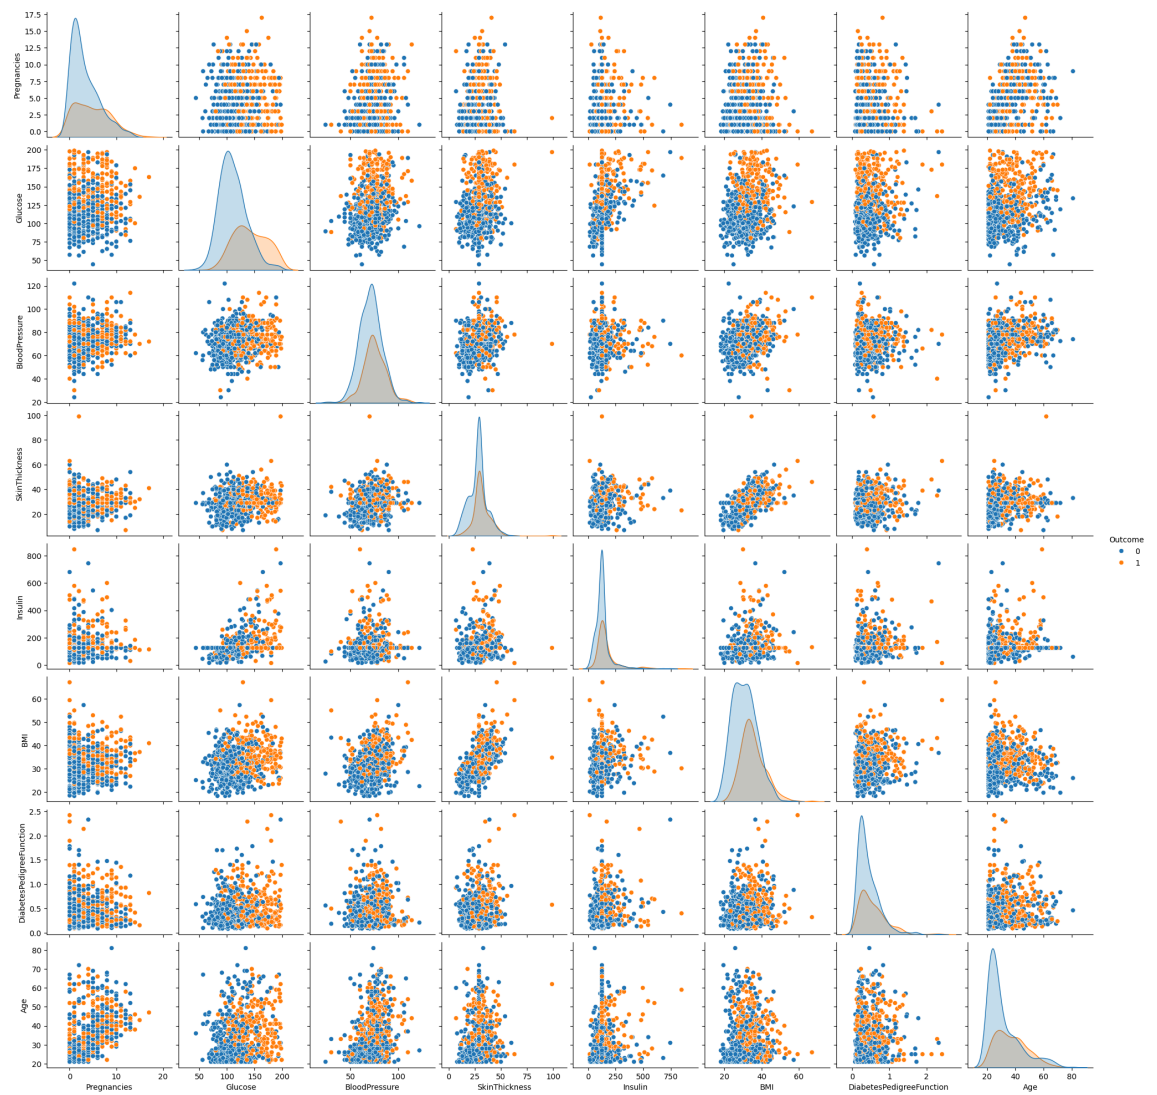


```
In [15]: 1 from pandas.plotting import scatter_matrix
```


```
In [16]: 1 p = scatter_matrix(df,figsize=(25,25))
```



```
In [17]: 1 p = sns.pairplot(df_copy, hue='Outcome')
```



```
In [18]: 1 plt.figure(figsize=(12,10))
2 p = sns.heatmap(df.corr(),annot=True,cmap='RdYlGn')
```



```
In [19]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [20]: 1 scale_X = StandardScaler()
```

```
In [21]: 1 X = scale_X.fit_transform(df_copy.drop(['Outcome'],axis=1),)
2 X = pd.DataFrame(X,columns=['Pregnancies','Glucose','BloodPressure','Sk
```

```
In [22]: 1 X.head()
```

```
Out[22]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPediq
0	0.639947	0.865108	-0.033518	0.670643	-0.181541	0.166619	
1	-0.844885	-1.206162	-0.529859	-0.012301	-0.181541	-0.852200	
2	1.233880	2.015813	-0.695306	-0.012301	-0.181541	-1.332500	
3	-0.844885	-1.074652	-0.529859	-0.695245	-0.540642	-0.633881	
4	-1.141852	0.503458	-2.680669	0.670643	0.316566	1.549303	



```
In [23]: 1 from sklearn.model_selection import train_test_split
```

```
In [24]: 1 y = df_copy.Outcome
```

```
In [25]: 1 X_train,X_test,Y_train,Y_test = train_test_split(X,y,test_size=1/3,rand
```

```
In [26]: 1 #K-Nearest-Neighbors
2
3 from sklearn.metrics import accuracy_score
4 from sklearn.neighbors import KNeighborsClassifier
5 testing_score=[]
6 training_score=[]
7 for i in range(1,15):
8     knn = KNeighborsClassifier(i)
9     knn.fit(X_train,Y_train)
10    training_score.append(knn.score(X_train,Y_train))
11    testing_score.append(knn.score(X_test,Y_test))
```

```
In [27]: 1 max_training_score = max(training_score)
2 train_scores_ind = [i for i,v in enumerate(training_score) if v == max_
```


```
In [28]: 1 print('Max training score{}% and k ={}'.format(max_training_score*100,1
2
```

Max training score100.0% and k =[1]

```
In [29]: 1 max_testing_score = max(testing_score)
2 test_scores_ind = [i for i,v in enumerate(testing_score) if v == max_te
3 print('Max training score{}% and k ={}'.format(max_testing_score*100,li
```

Max training score76.5625% and k =[11]

```
In [33]: 1 plt.figure(figsize=(12,5))
2
3 pplot = sns.lineplot(data=training_score,marker='*',label='Train Score')
4 pplot = sns.lineplot(data=testing_score,marker='o',label='Test Score')
```



```
In [34]: 1 knn = KNeighborsClassifier(11)
          2 knn.fit(X_train,Y_train)
          3 knn.score(X_test,Y_test)
```

Out[34]: 0.765625

```
In [ ]: 1
```